

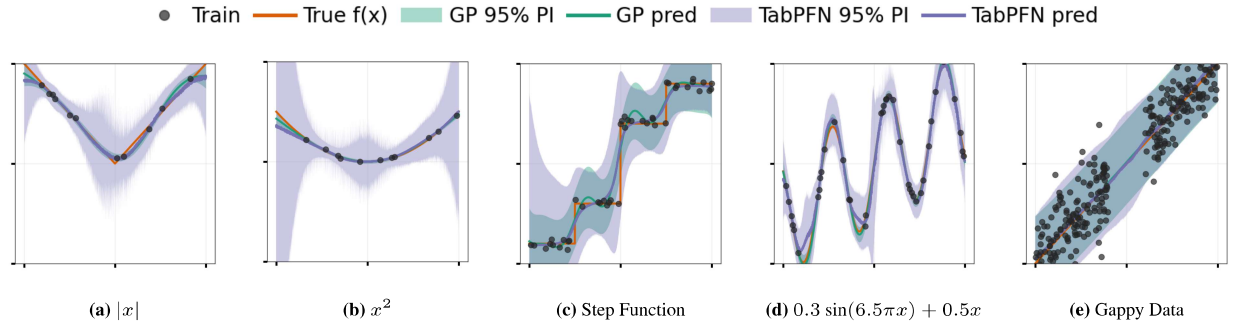


Figure 1 1D Examples: The functions are all adopted from [1]. Each plot has the same x and y axes, both spanning $[-0.5, 0.5]$. All the GPs are trained on CPU using the defaults of [2] (e.g., constant mean and a Gaussian kernel). TabPFN v2.5 is run on GPU without fine-tuning.